JORDAN A. RIVERA

(555) 014-2288 | jordan.rivera@example.com | linkedin.com/in/jordan-rivera-example | github.com/jordan-rivera-example | Portland, OR

EDUCATION

Cascadia State University — M.S. Computer Science — Expected Dec 2026
Coursework: Deep Learning, Artificial Intelligence, Advanced Algorithms, Software Engineering

Westbrook Institute, Pune, India — B.Tech, Electrical & Electronics Engineering — Jun 2022

EXPERIENCE

Cascadia State University — Graduate Research Assistant — Oct 2026 – Present
Portland, OR
- Fine-tune open-source large language models for research on LLM agents
- Analyze agent trajectories and design guardrails that constrain agent tool use and behavior

Nordvale Consulting — Software Engineer — Feb 2024 – Jul 2025
Pune, India
Client: Autolane — retail platform, 500K+ monthly active users
- Built 15+ React.js screens for Autolane's vehicle checkout platform using Redux Toolkit and TypeScript, shipped across web and mobile
- Built 10+ production features (React.js, TypeScript, Node.js), cutting defect rates by 30% and maintaining 99.9% uptime across distributed microservices
- Built a 20+ component reusable React library backed by Node.js REST APIs, adopted team-wide and cutting feature delivery time by 25%
- Containerized microservices with Docker and Kubernetes, raising deploy frequency 40% with zero-downtime releases via GitHub Actions CI/CD
- Led 15+ code reviews that cut post-release bugs by 28%, and mentored 2 junior engineers through a legacy jQuery-to-React migration
- Improved the platform to WCAG 2.1 accessibility and full responsiveness, raising usability scores by 20%
- Built analytics dashboards with GA4 and GTM, eliminating 3 hrs/week of manual reporting

Brightfold Systems — Associate Software Engineer — Jul 2022 – Dec 2023
India
- Built 8 full-stack applications (React.js, Next.js, Node.js) with event-driven flows via RabbitMQ
- Added server-side rendering with Next.js, cutting TTFB by 45% and pushing Core Web Vitals above 90 on Lighthouse
- Built 12+ REST and GraphQL APIs with JWT auth, rate limiting and versioning — 99.5% uptime across 50K+ daily requests
- Wired RabbitMQ async messaging between 5 core services, cutting cascading failure incidents by 60% during peak traffic
- Set up Docker deployments and GitHub Actions CI/CD across staging and production, enabling zero-downtime releases

Pinecrest Labs — Software Engineer Intern — Oct 2021 – Apr 2022
India
- Built 5 full-stack applications (React.js, Node.js, Spring Boot) with Kafka-driven microservice communication, cutting user-reported bugs by 30%
- Sped up MySQL and MongoDB queries on 10K+ records by 40% through composite indexes and rewritten aggregation pipelines

PROJECTS

ORBIT — Autonomous Job Application Agent — Python, Playwright, PostgreSQL, Llama 3.3-70B, FastAPI
- Built an autonomous multi-stage agent covering job discovery, per-posting resume tailoring and application submission across 5 applicant tracking systems
- Designed a plan-then-inject rewriting architecture in which the model emits edit coordinates as JSON and Python applies them, eliminating the 18–42% content loss of direct LLM editing
- Built a crawler over 1,579 company job boards that screens 39,000+ postings in under 60 seconds using title, location and eligibility filters before any model call
- Grounded every generated resume claim in a verified-fact graph with DeBERTa NLI entailment checks, so no bullet can assert anything the candidate's record does not support

DentScan — Vehicle Damage Detection — PyTorch, YOLO11m, FastAPI, React.js
- Trained YOLO11m on the CarDD dataset using A100 GPUs on NAU's Summit HPC cluster, reaching mAP@0.5 = 0.700, a 41.1% improvement over prior SOTA (DCN+ at 0.496)
- Per-class results: glass shatter 0.950, tire flat 0.874, lamp broken 0.830

- First application of YOLO11m to CarDD, with Class-Balanced Loss, Grad-CAM explainability and SAM zero-shot segmentation; targeting IEEE T-ITS
- Shipped a full-stack demo (FastAPI + React.js) with real-time inference and bounding-box visualization

VERA-RAG — Self-Verifying Adaptive RAG System FAISS, BM25, DeBERTa-v3, Llama-3, vLLM

- Built a closed-loop RAG system with hybrid retrieval (FAISS + BM25 + Reciprocal Rank Fusion), improving retrieval precision 23% over single-pass baselines
- Added DeBERTa-v3 NLI self-verification that detects hallucinated or unsupported claims and triggers agentic re-retrieval and regeneration
- Served Llama-3-8B via vLLM with structured prompt templates and function-calling patterns

SKILLS

Languages: Python, Java, JavaScript, TypeScript, C, SQL
Frontend: React.js, Next.js, Redux Toolkit, Tailwind CSS, HTML5, CSS3
Backend: Node.js, Express.js, Java Spring Boot, FastAPI, Flask, REST APIs, GraphQL
AI/ML: PyTorch, YOLO, FAISS, DeBERTa, LangChain, Hugging Face, vLLM, OpenAI API
Databases: PostgreSQL, MySQL, MongoDB, Redis
Cloud & Infra: Docker, Kubernetes, AWS (EC2, S3, Lambda), GitHub Actions CI/CD, SLURM
Tools: Git, Playwright, Postman, Jira